

Lecture 13:


Memory Management

John Wickerson


Compilers

Memory management


static



stack



heap

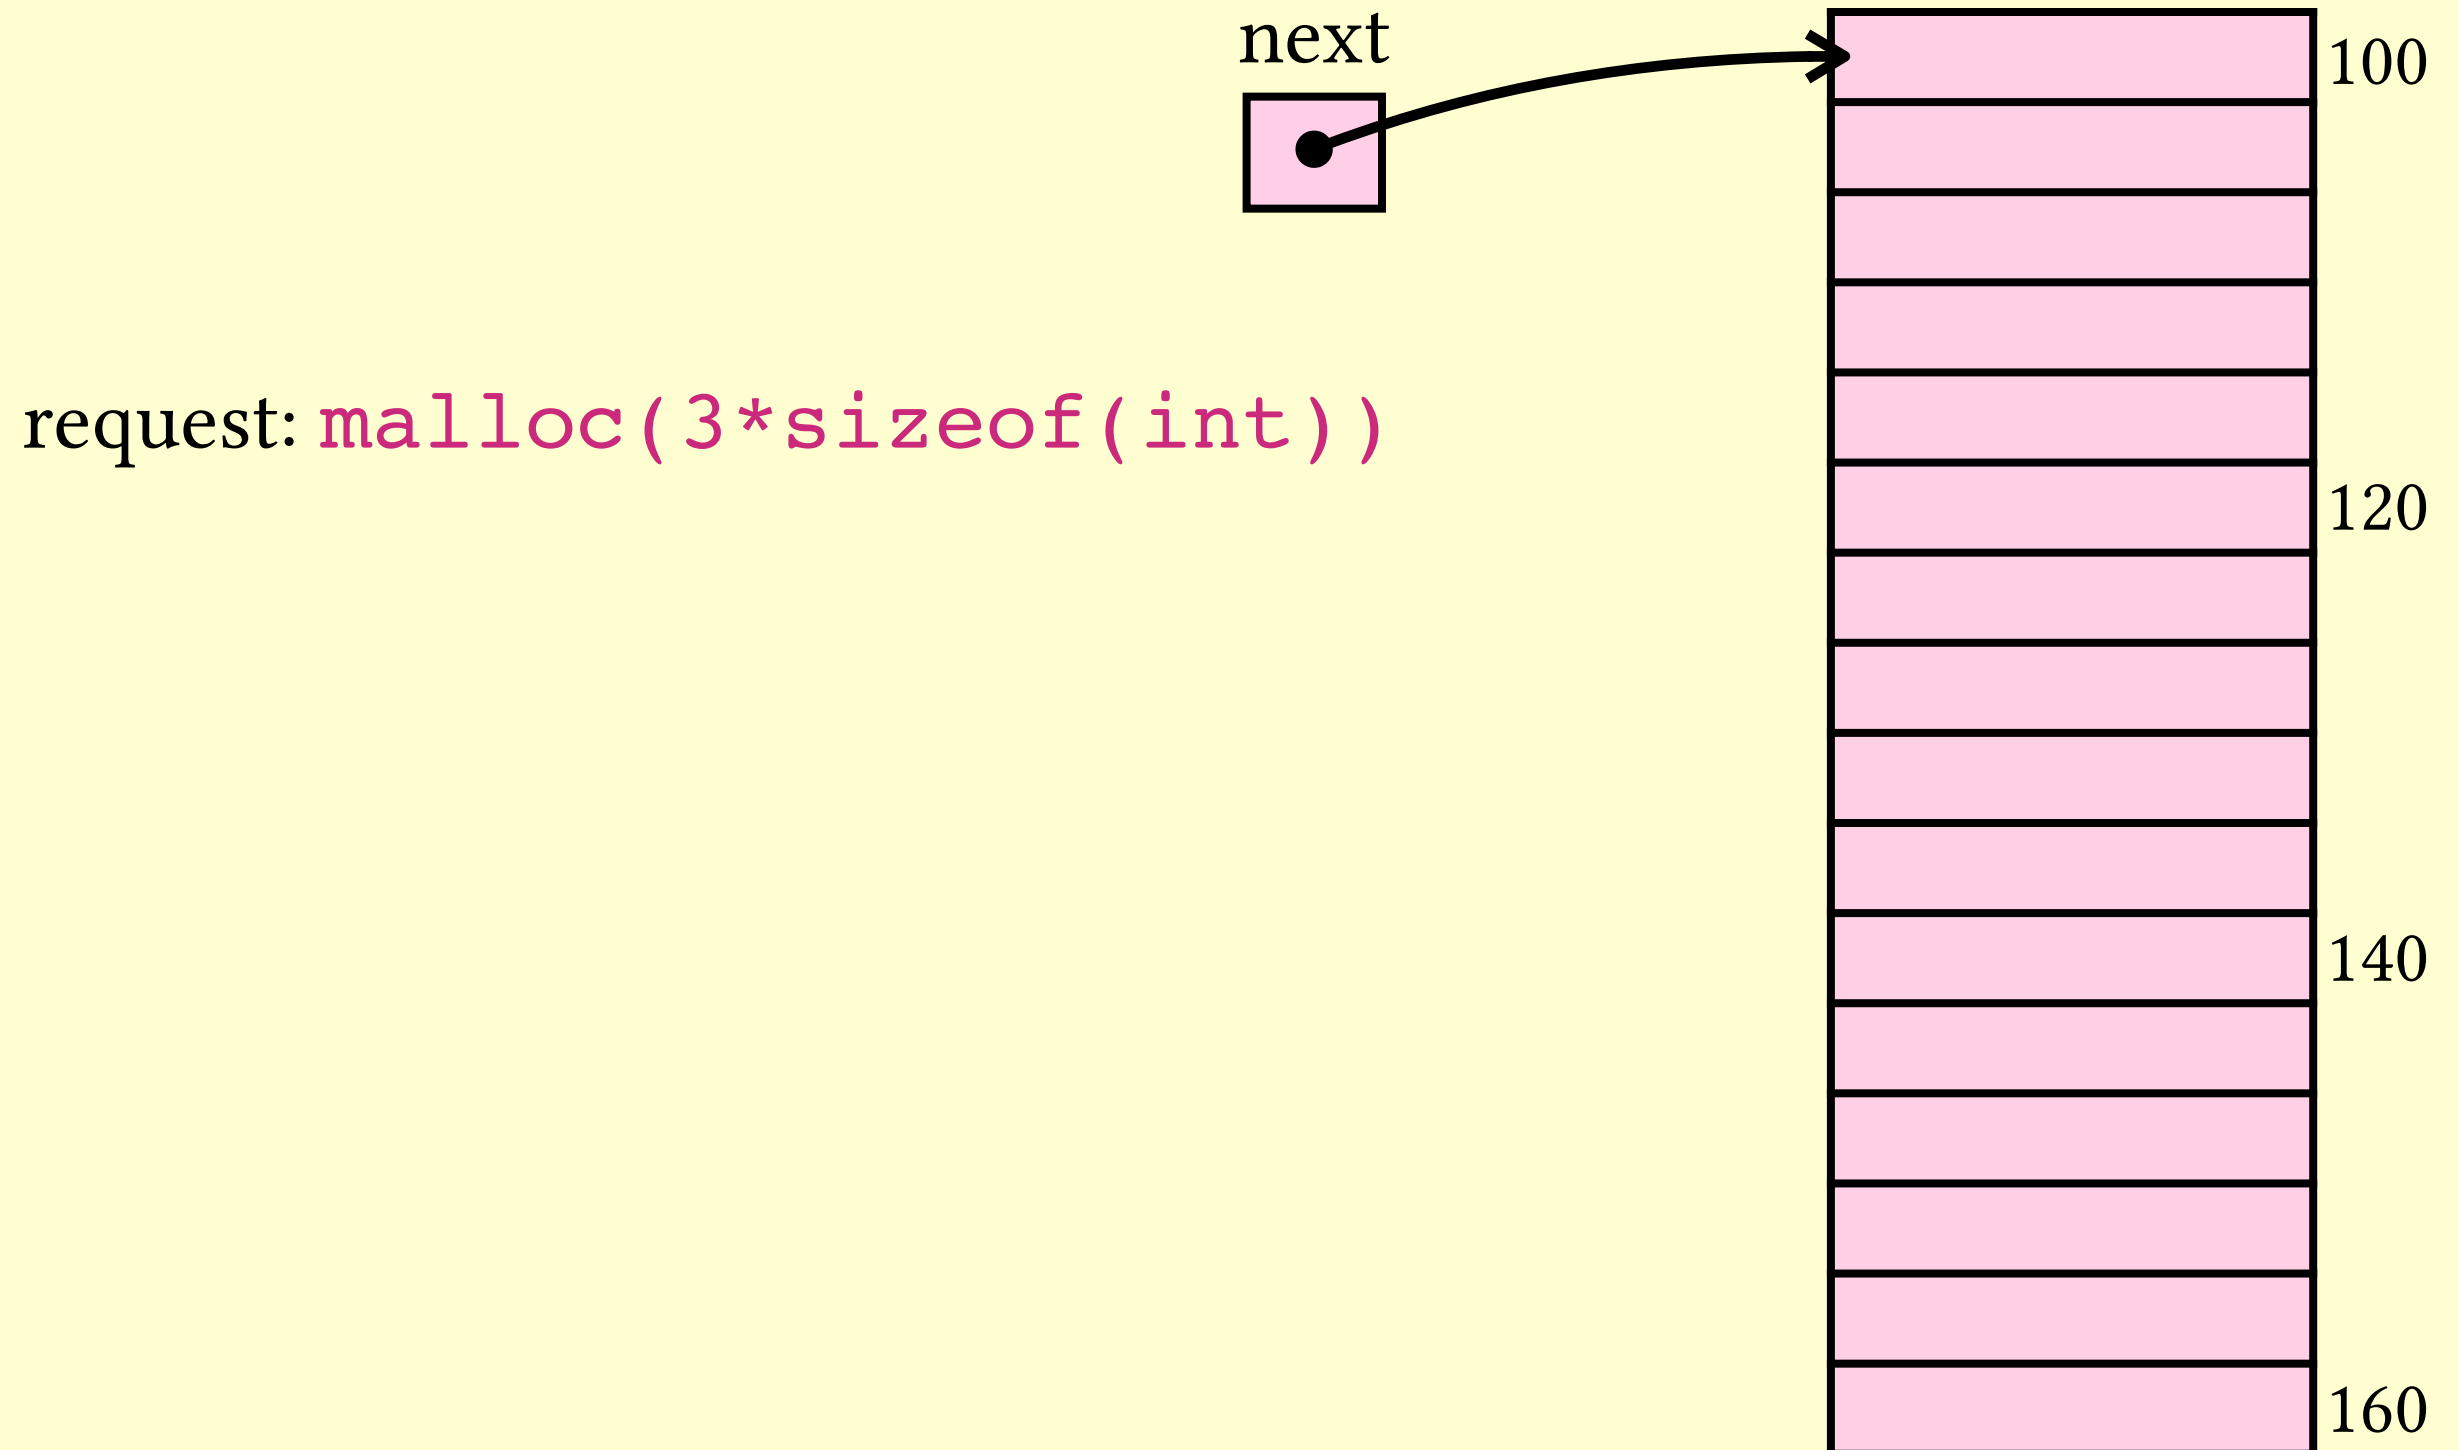


- Static allocation (e.g. global variables).
- Automatic/stack allocation (e.g., local variables).
- Dynamic/heap allocation, managed by programmer (i.e., `malloc`, `free`).
- Dynamic/heap allocation, managed automatically (i.e., garbage collection).

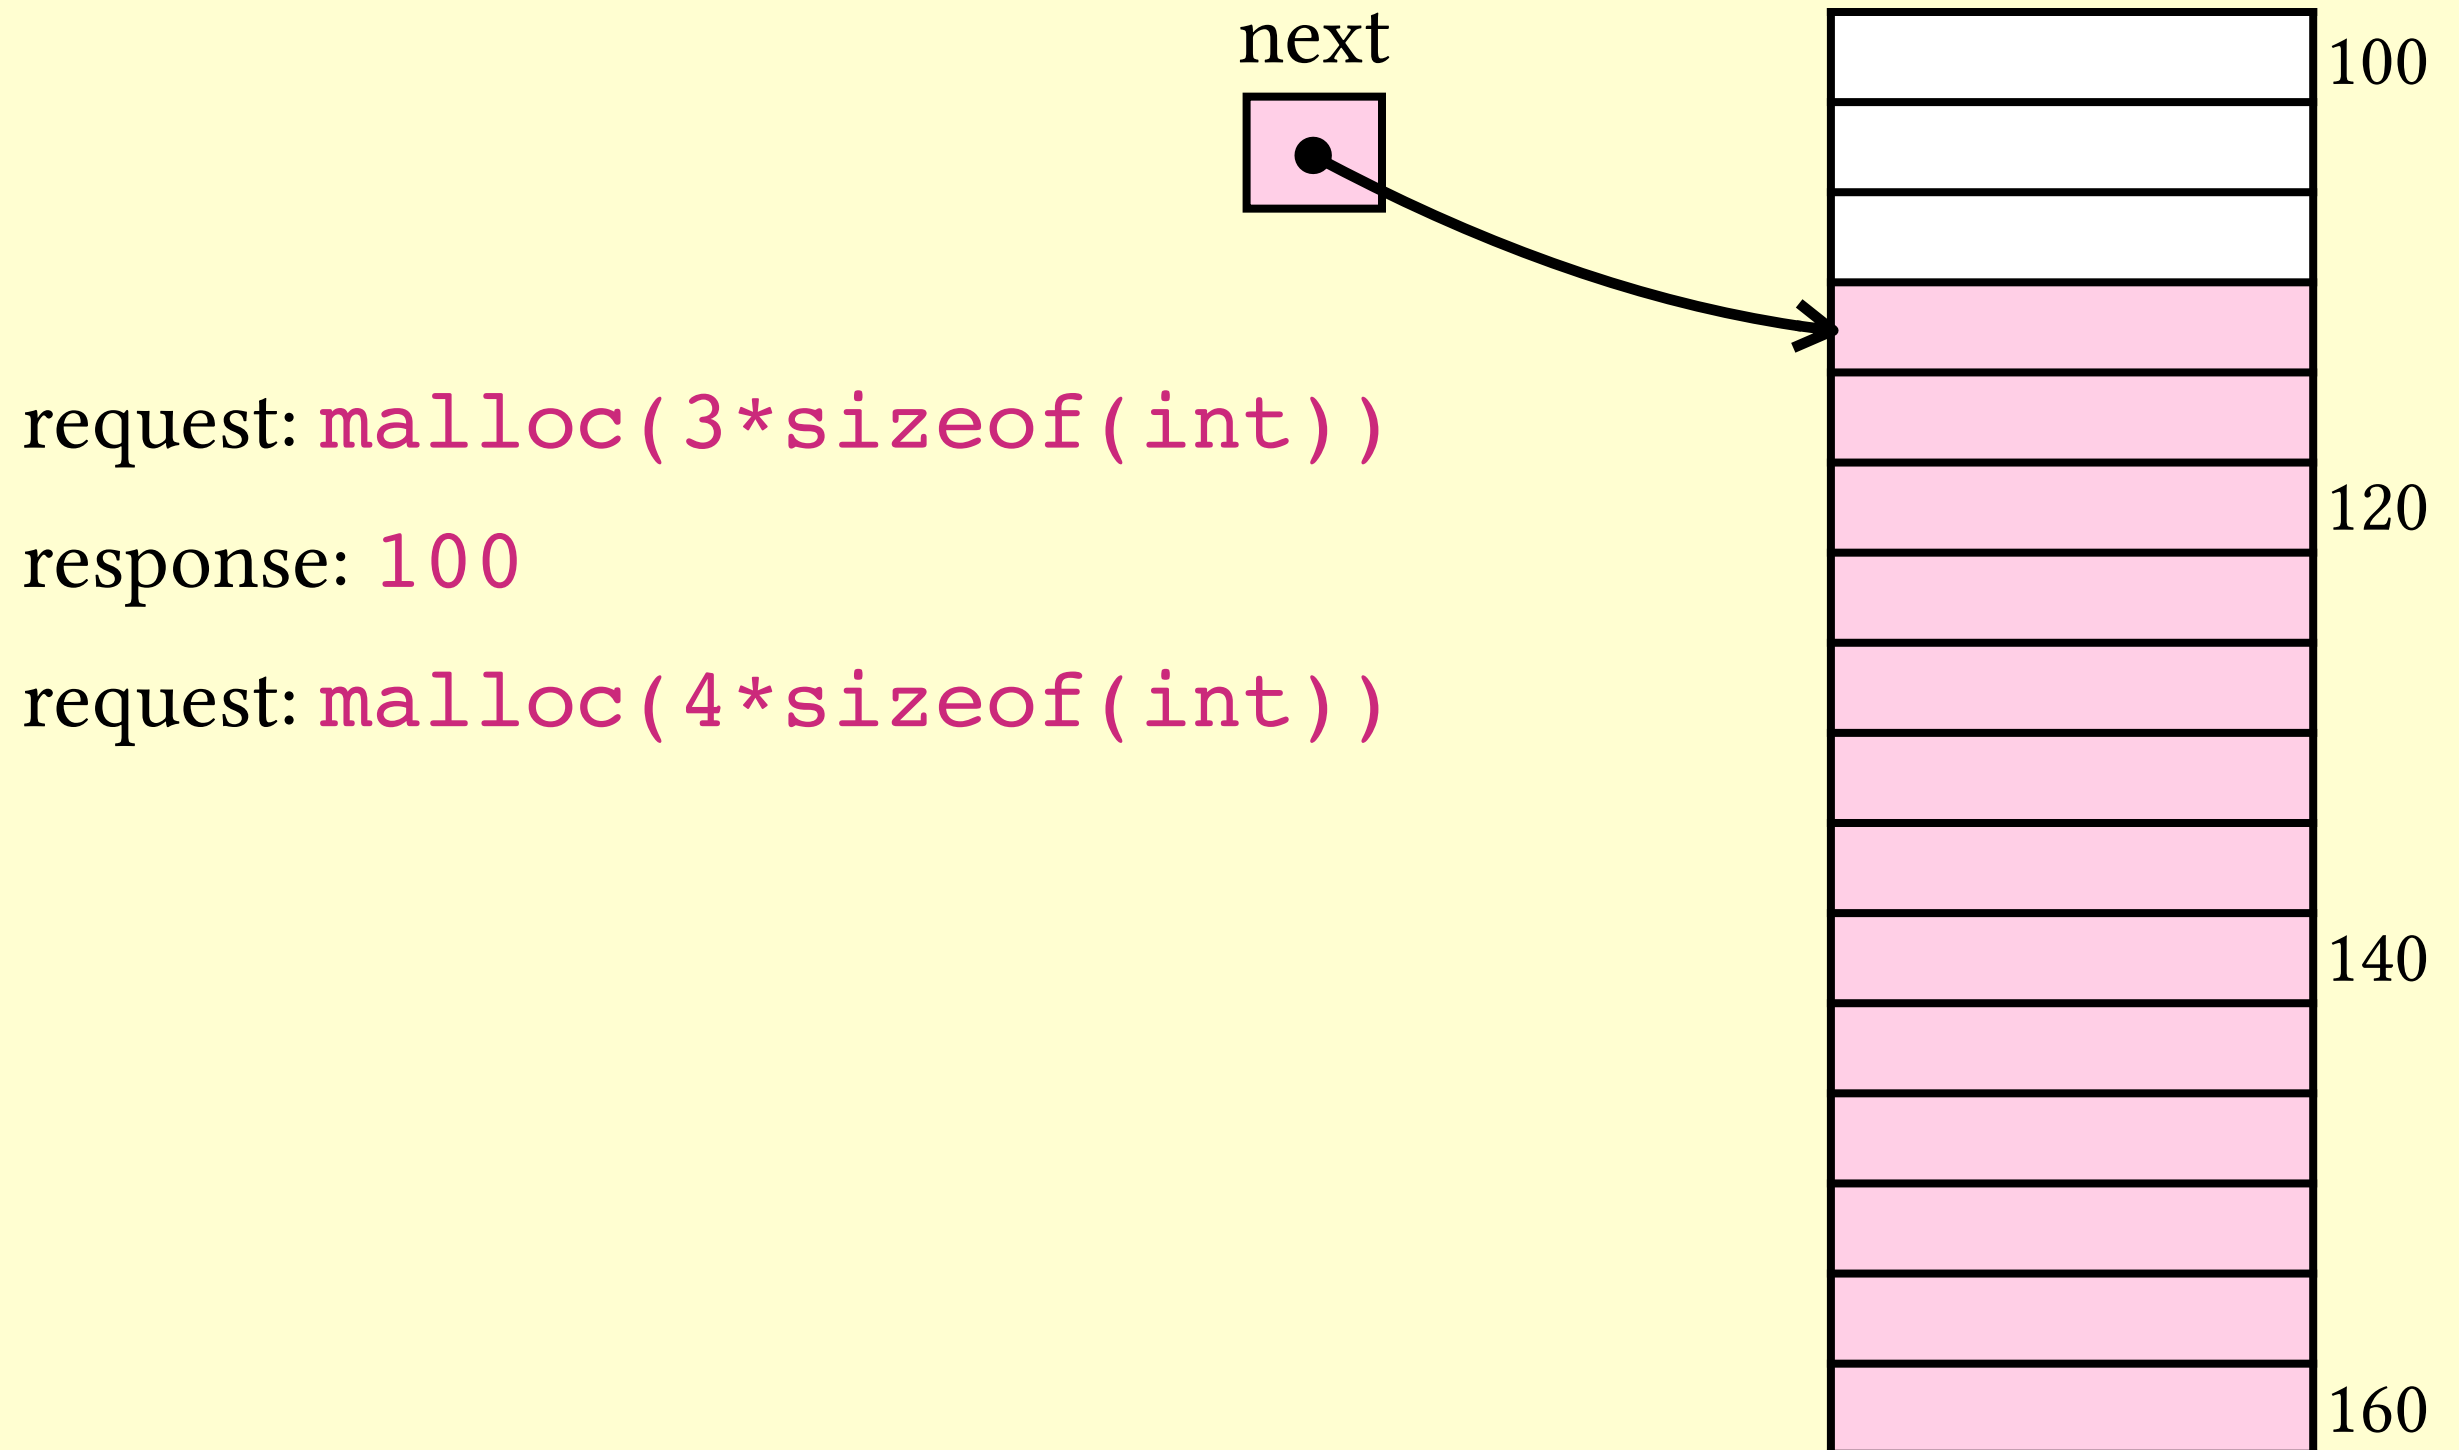
Lecture outline

- Dynamic memory allocation (`malloc` and `free`)
- Garbage collection

First malloc



First malloc



First malloc

request: `malloc(3*sizeof(int))`


response: `100`

request: `malloc(4*sizeof(int))`

response: `112`

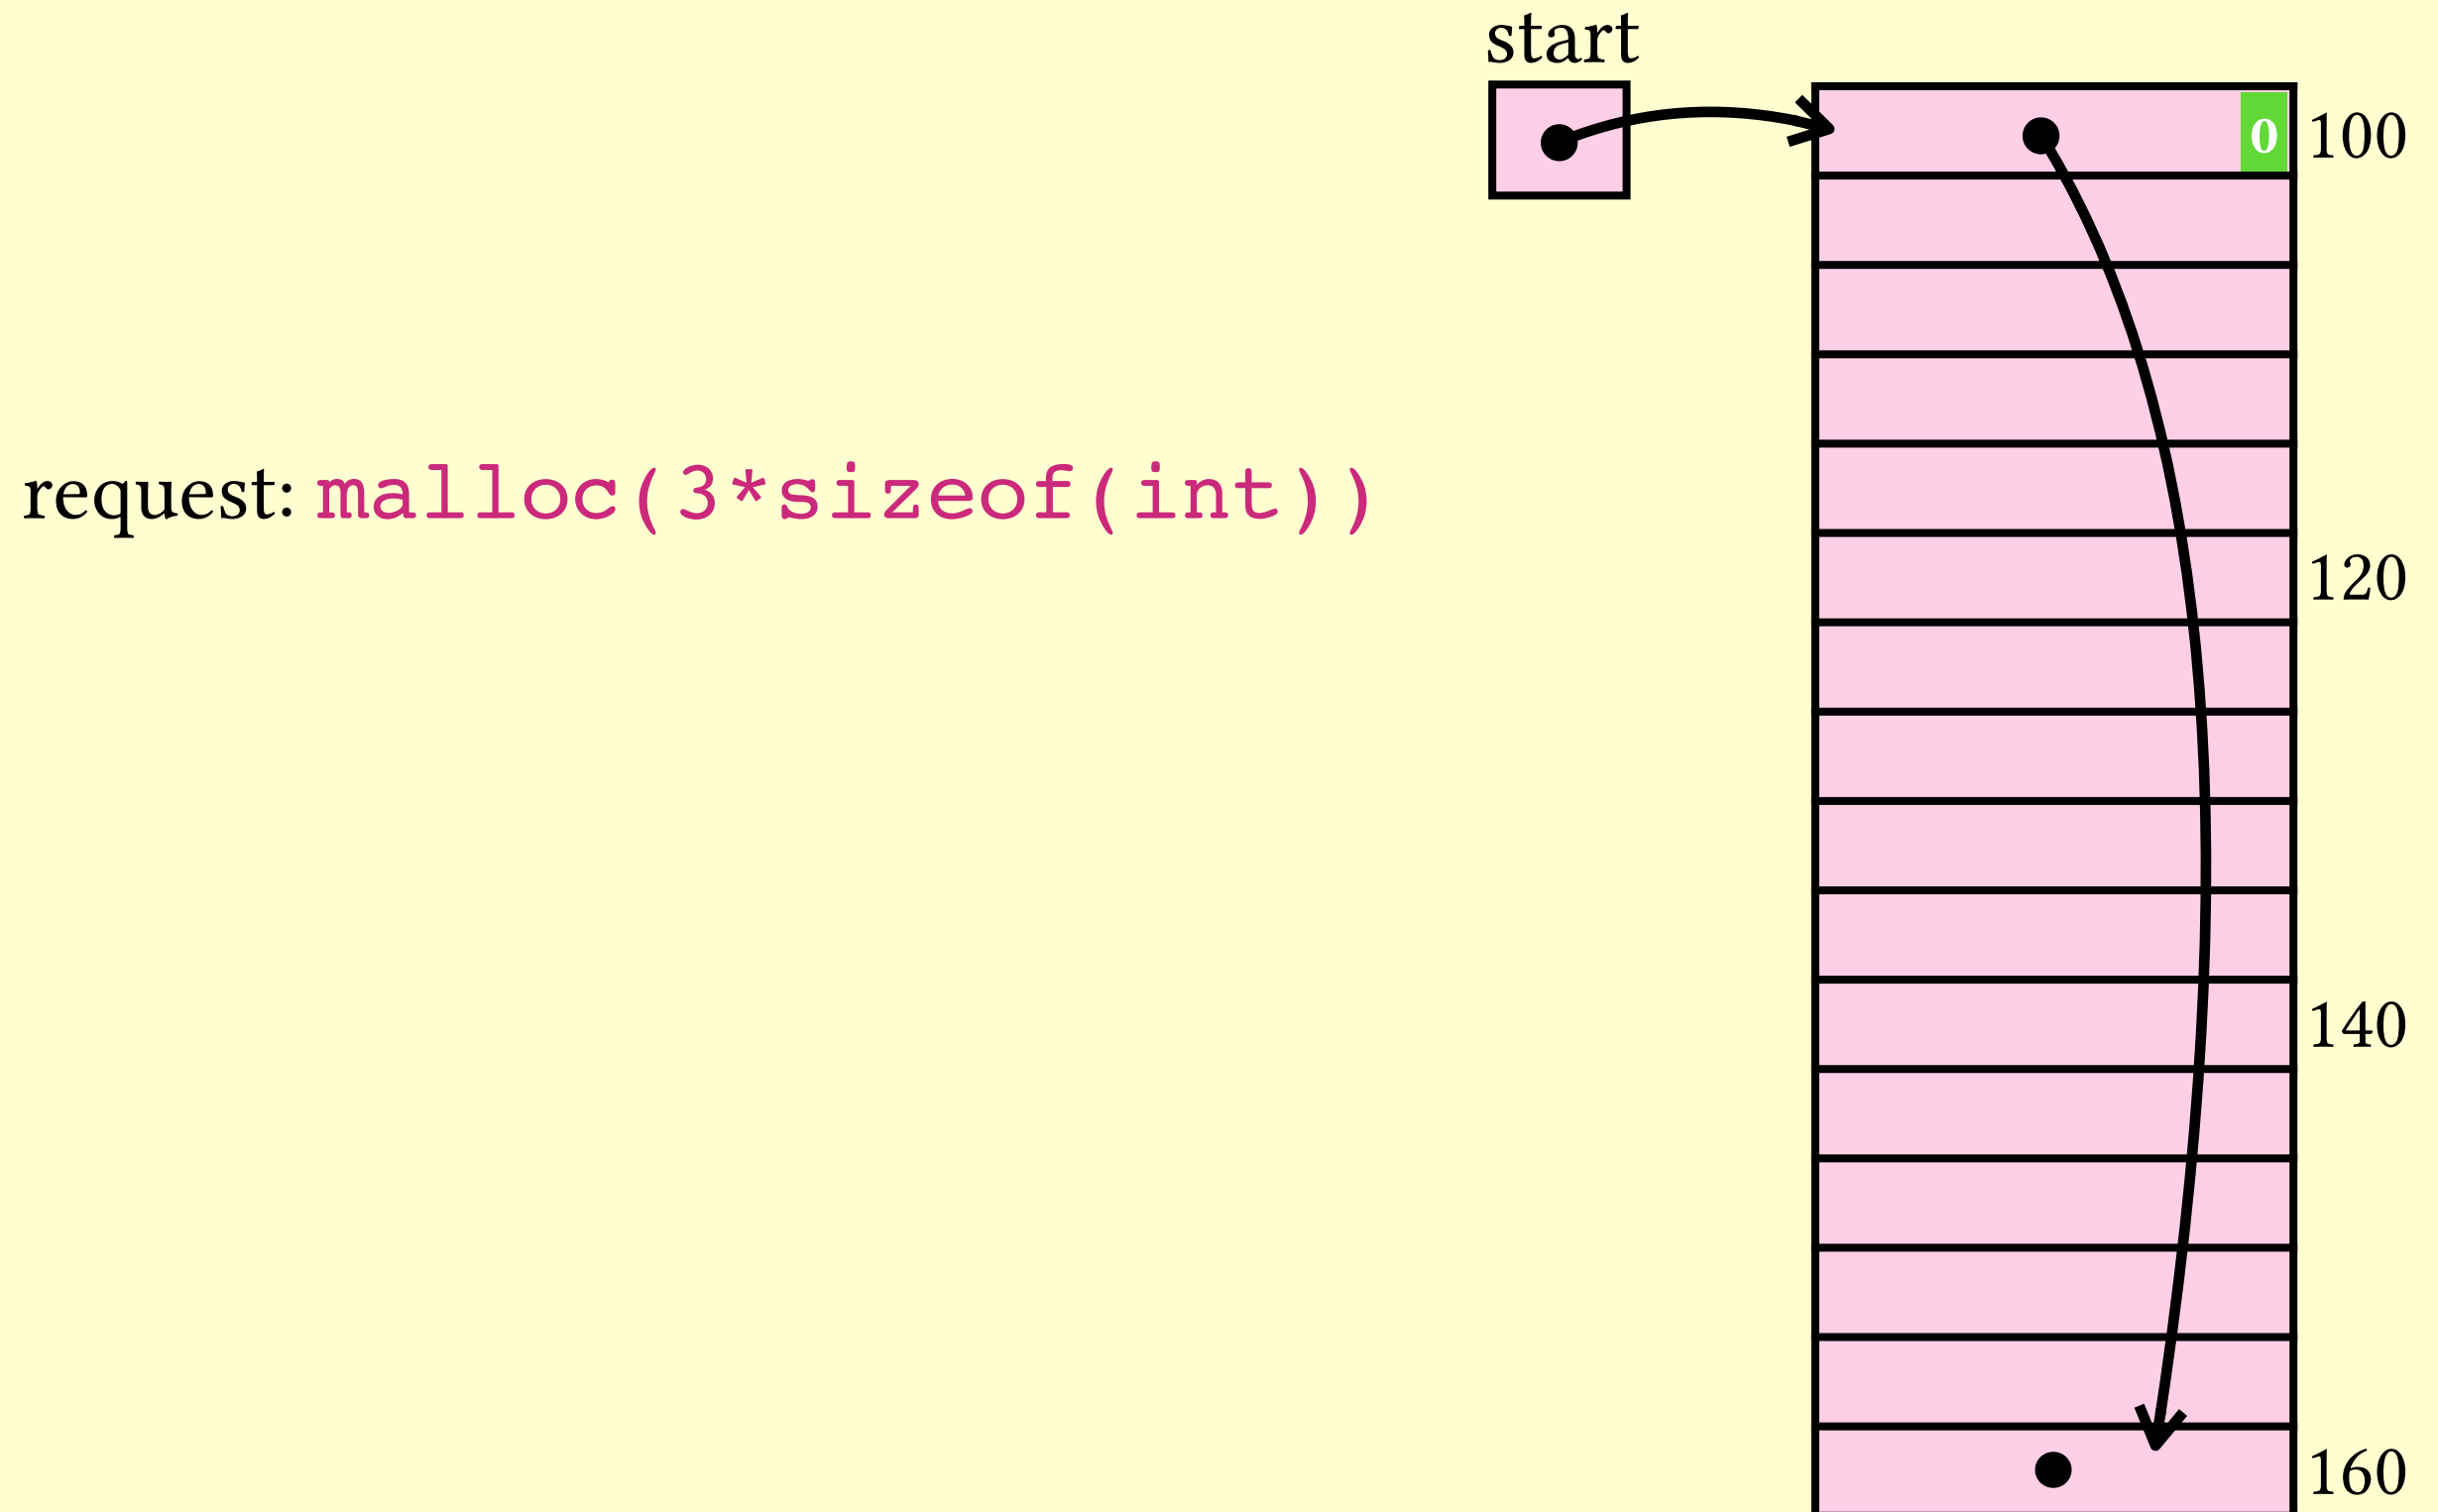
request: `free(100)`

response: `[done]`



Can we do better?

Second malloc




Second malloc

request: `malloc(3*sizeof(int))`

response: 104

request: `malloc(4*sizeof(int))`




Second malloc

request: `malloc(3*sizeof(int))`

response: 104

request: `malloc(4*sizeof(int))`

response: 120





Second malloc

request: `malloc(3*sizeof(int))`

response: 104

request: `malloc(4*sizeof(int))`

response: 120



Second malloc


request: `malloc(3*sizeof(int))`

response: 104

request: `malloc(4*sizeof(int))`

response: 120

request: `free(104)`



Second malloc

request: `malloc(3*sizeof(int))`

response: 104


request: `malloc(4*sizeof(int))`

response: 120

request: `free(104)`

response: [done]

request: `malloc(1*sizeof(int))`



Second malloc

request: `malloc(3*sizeof(int))`

response: 104

request: `malloc(4*sizeof(int))`


response: 120

request: `free(104)`

response: [done]

request: `malloc(1*sizeof(int))`

response: 104



Second malloc

request: `malloc(3*sizeof(int))`

response: 104

request: `malloc(4*sizeof(int))`


response: 120

request: `free(104)`

response: [done]

request: `malloc(1*sizeof(int))`

response: 104



request: `malloc(3*sizeof(int))`

response: 104 **Second malloc**

request: `malloc(4*sizeof(int))`

response: 120


request: `free(104)`

response: [done]

request: `malloc(1*sizeof(int))`

response: 104

request: `malloc(5*sizeof(int))`



request: `malloc(3*sizeof(int))`

response: 104 **Second malloc**

request: `malloc(4*sizeof(int))`

response: 120

request: `free(104)`

response: [done]


request: `malloc(1*sizeof(int))`

response: 104

request: `malloc(5*sizeof(int))`

response: 140

request: `free(120)`



request: `malloc(3*sizeof(int))`

response: 104 **Second malloc**

request: `malloc(4*sizeof(int))`

response: 120

request: `free(104)`

response: [done]

request: `malloc(1*sizeof(int))`


response: 104

request: `malloc(5*sizeof(int))`

response: 140

request: `free(120)`

response: [done]



request: `malloc(3*sizeof(int))`

response: 104 **Second malloc**

request: `malloc(4*sizeof(int))`

response: 120

request: `free(104)`

response: [done]

request: `malloc(1*sizeof(int))`


response: 104

request: `malloc(5*sizeof(int))`

response: 140

request: `free(120)`

response: [done]



Second malloc

Explicit Stabilisation for Modular Rely-Guarantee Reasoning

John Wickerson, Mike Dodds and Matthew Parkinson

University of Cambridge Computer Laboratory

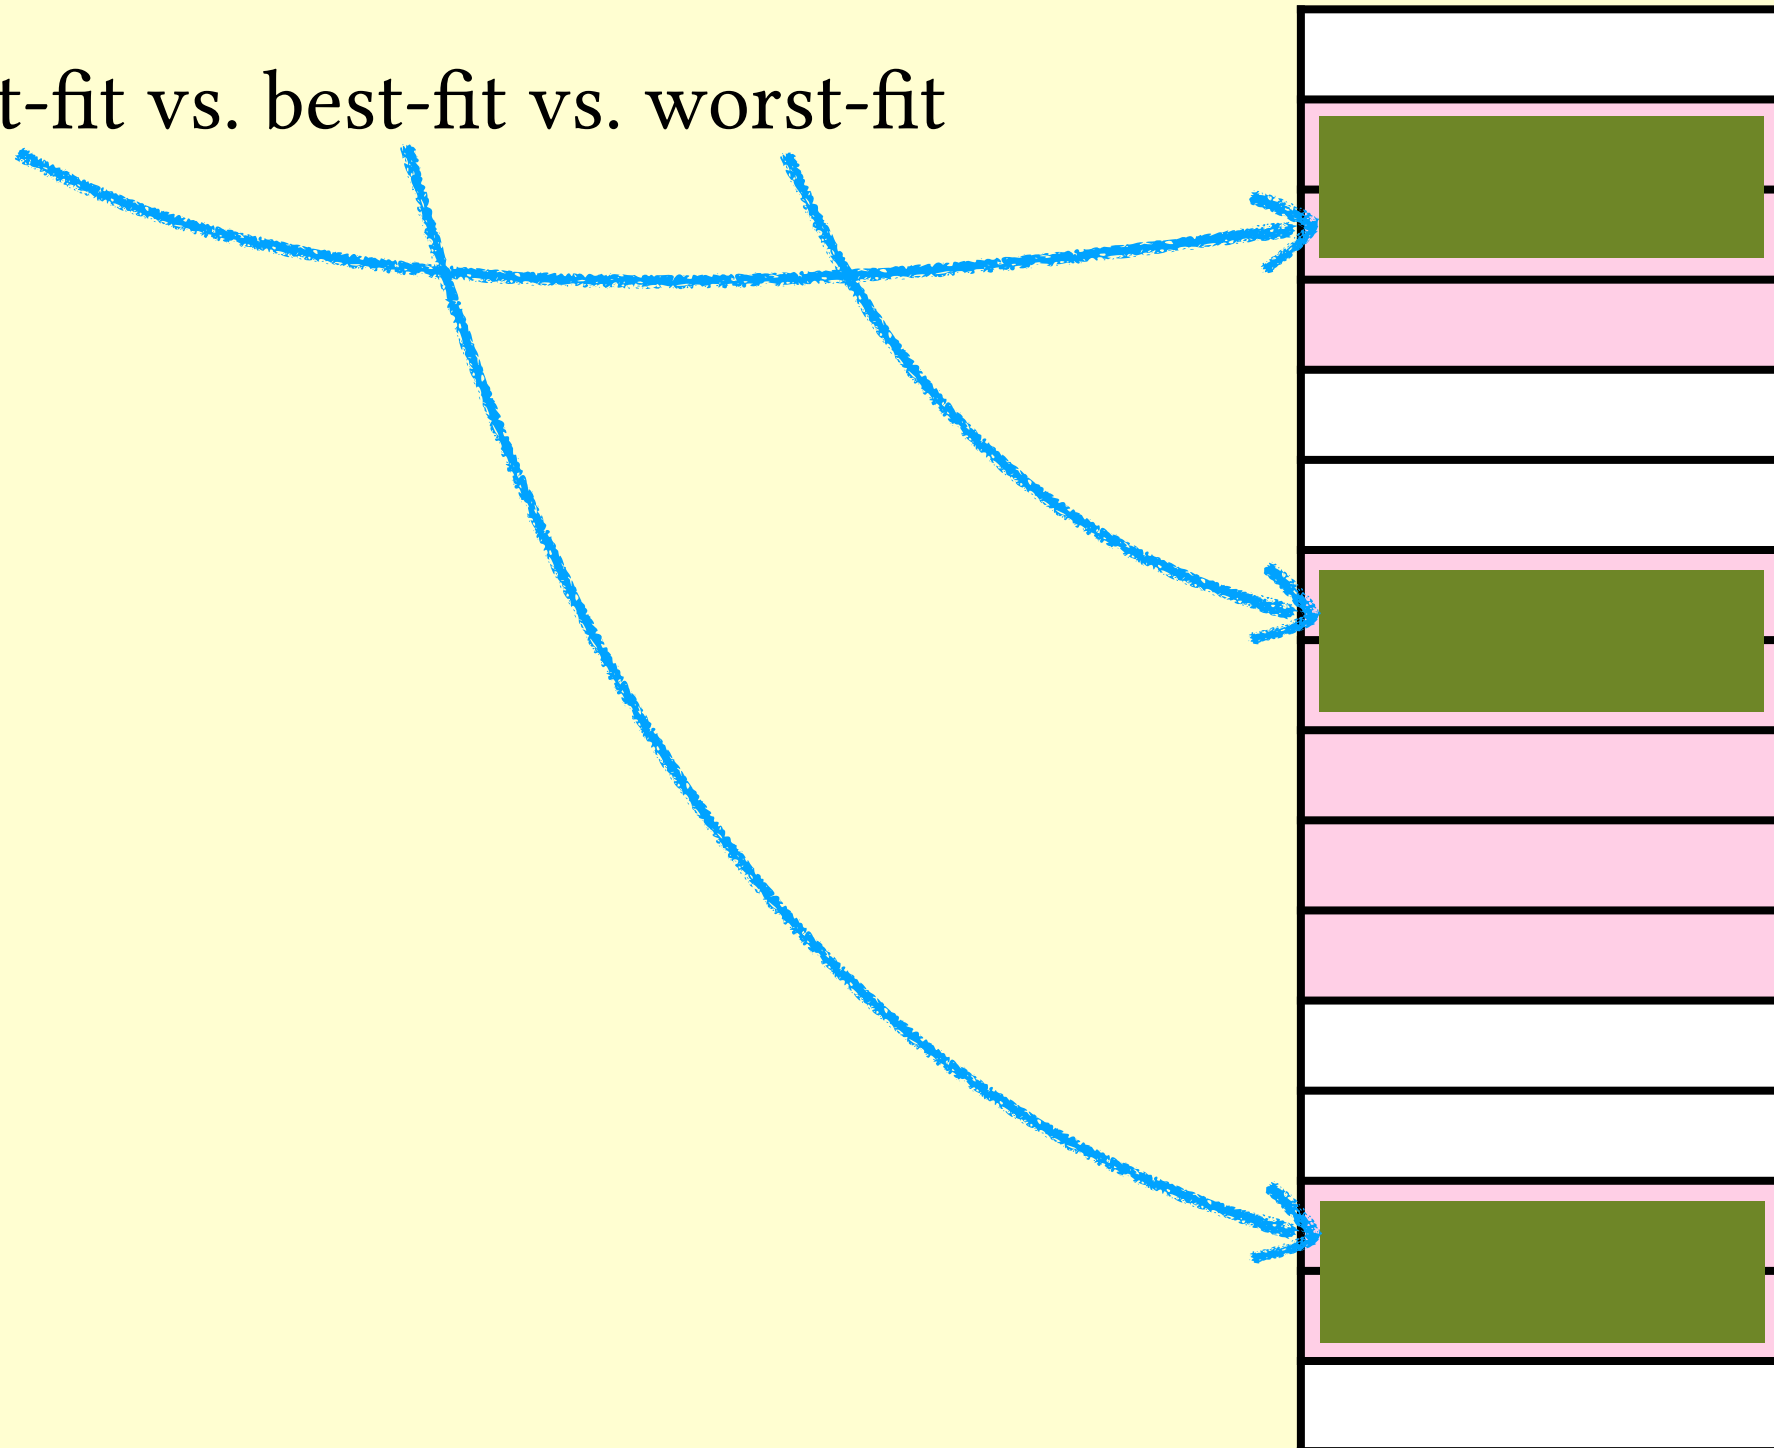
Abstract. We propose a new formalisation of stability for Rely-Guarantee, in which an assertion's stability is encoded into its syntactic form. This allows two advances in modular reasoning. Firstly, it enables Rely-Guarantee, for the first time, to verify concurrent libraries independently of their clients' environments. Secondly, in a sequential setting, it allows a module's internal interference to be hidden while verifying its clients. We demonstrate our approach by verifying, using RGSep, the Version 7 Unix memory manager, uncovering a twenty-year-old bug in the process.

1 Introduction

Verifying concurrent programs is hard because commands from different


Malloc decisions

- First-fit vs. best-fit vs. worst-fit



Malloc decisions

- First-fit vs. best-fit vs. worst-fit
- **Exercise.** Devise a sequence of `malloc` calls where worst-fit is the best of the three strategies.




Can we do better?


Third malloc

- Doug Lea's malloc (**dlmalloc**) was invented around 1987. Its derivatives are widely used in GNU/Linux.



Doug Lea





Can we do better?

Other mallocs


- `phkmalloc`, `ptmalloc`, `jemalloc`, `snmalloc`, ...
- Better performance when multiple threads are concurrently calling `malloc` and `free`.

Lecture outline

- ✓ Dynamic memory allocation (`malloc` and `free`)
- Garbage collection


Garbage collection

- Invented around 1959 by John McCarthy and implemented in his LISP programming language.
- Alternative to **free**. The garbage collector automatically identifies memory that will never be used again by the program, and makes it available for reallocation.
- **Question.** Which languages use garbage collection?
- **Question.** What are the advantages/disadvantages of GC?




Reference counting


static




stack



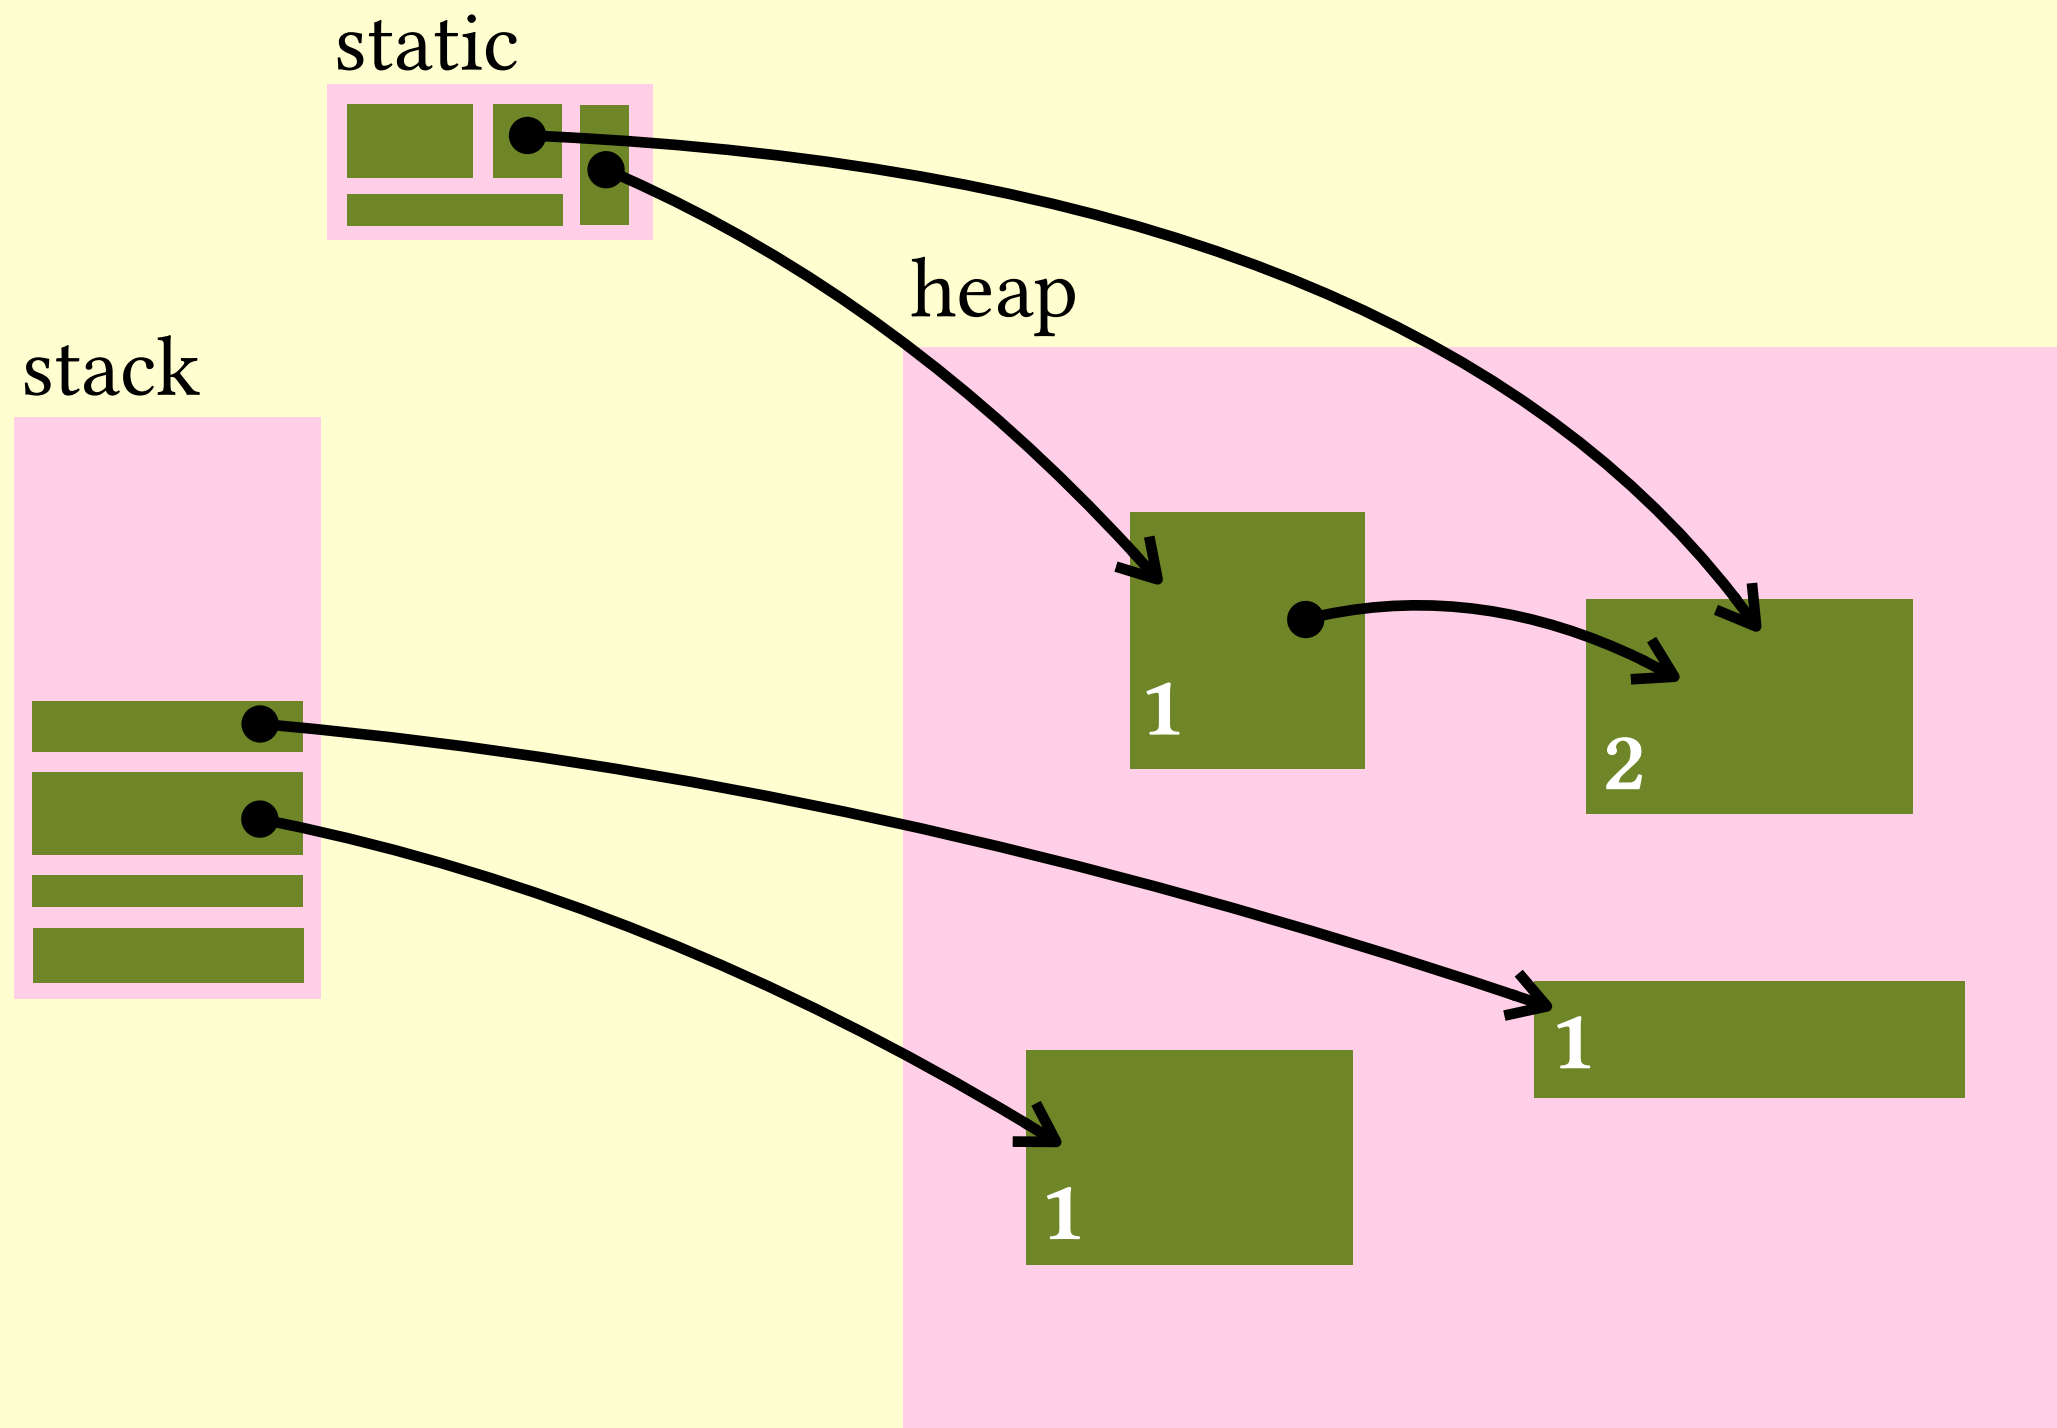
heap




Reference counting




Reference counting




Reference counting




Reference counting



Tracing GC



Tracing GC



Comparison

- **Question.** What are the advantages/disadvantages of Reference counting vs. Tracing GC?

- Common problem: pointer arithmetic.

```
int *p = malloc(sizeof(int));  
p += 1000;
```

Do we still have a reference?

- Common problem: imprecision.

```
int *p = malloc(sizeof(int));  
*p = 42;  
int *q = malloc(sizeof(int));  
return;
```


*Ref. counting will reclaim **p** here*

*Tracing GC will reclaim **p** here*

*But we could reclaim **p** here*


Other GCs

- Copying GC




Other GCs

- Copying GC




Other GCs

- Generational GC



Other GCs

- Generational GC



Lecture outline

- ✓ Dynamic memory allocation (`malloc` and `free`)
- ✓ Garbage collection

Summary

- Memory regions: static, stack, heap.
- Memory allocation should: be fast, be responsive, avoid fragmentation, be aware of spatial locality, use little memory, and be concurrency-safe.
- Heap deallocation is manual (**free**) or automatic (GC).
- GC should: be fast, be responsive, be aware of spatial locality, use little memory, be concurrency-safe, be sound (don't collect non-garbage), and be precise (collect all garbage).